Founding Product Engineer – Software

London (Hybrid) · Reports to Head of Labs

THE OPPORTUNITY

An early founding hire inside a fast-moving internal tech lab building two core products: **Mind** (external-facing product for hiring managers) and **Mothership** (internal data lake, copilot, and dashboard powering the operating system and feeding Mind).
You'll own the product surfaces across both — building intuitive hiring workflows for external users in Mind, and internal copilot and dashboard tooling in Mothership. This isn't a narrow front-end role. You'll work across UI, APIs, integrations, and deployment end-to-end.

WHAT YOU'LL DO

- Build task-based, workflow-first interfaces for hiring managers using React/Next.js
- Build copilot-style dashboards and internal tooling that turns structured data into useful actions
- Turn manual internal workflows into repeatable, productised experiences
- Own delivery end-to-end — UI through to deployment
- Build against shared data contracts and canonical object definitions
- Make pragmatic architectural decisions that allow fast iteration without sacrificing maintainability
- Build AI-assisted workflows with guardrails, transparency, and user control
- Instrument adoption, deploy regularly, and iterate on real usage data
- Understand how product decisions affect activation, retention, and revenue

ABOUT YOU

- 4–7+ years building scalable web products — B2B SaaS, marketplaces, workflow tools, or talent tech
- Experience in a 0–1 or early-stage build — you know what shipping fast actually looks like
- Strong full-stack: React/Next.js, Node.js, TypeScript, Python where needed, PostgreSQL
- Product-minded — UX and outcomes matter as much as code quality
- Comfortable turning ambiguity into a clear build plan and shipping a pragmatic v1
- Data-literate: APIs, analytics, data-heavy workflows

CURRENT STACK DIRECTION React/Next.js · Node.js / TypeScript · Python · PostgreSQL · Supabase

Founding Product Engineer – Data

London (Hybrid) · Reports to Head of Labs

THE OPPORTUNITY

Same lab, different layer. This role is about building the brain behind both products. Mothership is the foundation — simultaneously a data lake, a copilot layer, and a dashboard product. You'll build what makes it defensible: the data model, pipelines, retrieval layer, quality controls, and data contracts everything else runs on.

Competitors can copy a UI. They can't easily replicate structured operational signal, real-world feedback loops, and workflow context turned into reliable product intelligence.

WHAT YOU'LL DO

- Design and maintain the centralised data platform ingesting from operational systems (e.g. Bullhorn)
- Build canonical data models for candidates, roles, clients, processes, interactions, and outcomes — including identity resolution and deduplication
- Build ETL/ELT pipelines with monitoring, data quality checks, and production reliability
- Define and maintain data contracts and shared object definitions that product surfaces build against
- Build the signal layer powering Mothership copilots and dashboards: event tracking, action triggers, recommendation primitives, and decision-grade visibility
- Integrate LLM capabilities responsibly — extraction, ranking, contextual triggers — with traceability and confidence signalling
- Build analytics foundations to understand product usage, workflow impact, and outcome drivers
- Optimise schemas, indexing, caching, and update patterns for performance and scale

ABOUT YOU

- 4–8+ years building data-heavy platforms, backend systems, or workflow products
- You've taken messy real-world data and turned it into reliable, structured systems
- Strong Python and PostgreSQL; experience with pipeline design, schema design, APIs, monitoring, and third-party integrations
- Data product mindset — you think in objects, contracts, retrieval, and usability, not just storage
- LLM/embeddings experience useful, but reliability and product usefulness matter more than novelty
- Can make architectural calls under ambiguity and keep standards high while moving quickly

CURRENT STACK DIRECTION Python · PostgreSQL · ETL/ELT Pipelines · APIs · Indexing / Search · Embeddings where required